

Proyecto de Grado: IA al servicio de las comunidades

Maestría en Analítica para la Inteligencia de Negocios

Nombres: María Fernanda Izquierdo Aparicio & Ana María Ochoa Muñoz

Fecha: 18 de Abril de 2025

Este modelo forma parte del proyecto de grado de la Maestría en Analítica para la Inteligencia de Negocios de la Pontificia Universidad Javeriana, enfocado en la aplicación de Inteligencia Artificial al servicio de las comunidades. Su objetivo es clasificar imágenes de plantas en función de su especie usando una arquitectura pre-entrenada RESNET

1. Importe de Librerías

```
In [4]: import json

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import matplotlib.image as mpimg

import os
import zipfile
import shutil
import seaborn as sns
import datetime
import cv2

import keras
from keras import models
from keras import layers

import random

import tensorflow as tf
from tensorflow.keras.models import Model
from tensorflow.keras import Layer
from tensorflow.keras.layers import Input, Dense
from tensorflow.keras.preprocessing import image
```

```

from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.preprocessing.image import img_to_array, load_img
from tensorflow.keras.callbacks import ReduceLROnPlateau
from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping, ReduceLROnPlateau
from tensorflow.keras.applications import EfficientNetB0
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.applications import ResNet50
from tensorflow.keras.layers import GlobalAveragePooling2D
from tensorflow.keras.utils import to_categorical

from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, f1_score
from sklearn.metrics import confusion_matrix, accuracy_score, ConfusionMatrixDisplay
from sklearn.metrics import roc_curve, auc
from sklearn.metrics import classification_report

from PIL import Image # Librería para abrir y manipular imágenes

# ignoring warnings
import warnings
warnings.simplefilter("ignore")

```

```

In [5]: #Medir_tiempo_ejecucion():
        inicio = datetime.datetime.now()
        print(f"Inicio: {inicio.strftime('%Y-%m-%d %H:%M:%S.%f')}")

```

Inicio: 2025-04-18 18:44:26.469952

2. Definición de Rutas y Parámetros

```

In [7]: data_dir = 'Datos Finales' # Reemplaza con la ruta real
        metadata_file = 'Metadatos_Finales.csv' # Reemplaza con la ruta real
        image_height = 150
        image_width = 150
        batch_size = 15
        num_classes = None # Se determinará automáticamente del archivo CSV

```

3. Cargar y procesar los metadatos

```

In [9]: try:
        metadata = pd.read_csv(metadata_file, sep=";")
        metadata = metadata[metadata["Planta"] != "Orange"]
        metadata = metadata[metadata["Planta"] != "Grape"]
        metadata = metadata[metadata["Planta"] != "Pepper_bell"]
        metadata = metadata[metadata["Planta"] != "Peach"]
        metadata = metadata.rename(columns={'Nombre del archivo': 'image_id'})

```

```
# Asegúrate de que las columnas 'filename' (o el nombre de tu columna de
# y 'label' (o el nombre de tu columna de etiquetas) existan.

if 'image_id' not in metadata.columns or 'Planta' not in metadata.columns:
    raise ValueError("El archivo CSV debe contener las columnas 'image_id' y 'Planta'")
class_names = metadata['Planta'].unique().tolist()
num_classes = len(class_names)
print(f"Se encontraron {num_classes} clases: {class_names}")
except FileNotFoundError:
    print(f"Error: No se encontró el archivo de metadatos en '{metadata_file}'")
    exit()
except ValueError as e:
    print(f"Error al procesar el archivo CSV: {e}")
    exit()
```

Se encontraron 9 clases: ['Yuca', 'Strawberry', 'Potato', 'Corn_(maize)', 'Tomato', 'Apple', 'cauliflower', 'cucumber', 'bean']

```
In [10]: train, test = train_test_split(metadata, test_size=0.2, random_state=42, shuffle=True)
train.shape, test.shape
```

```
Out[10]: ((42354, 7), (10589, 7))
```

4. Crear un generador de datos

```
In [12]: datagen = ImageDataGenerator(
    validation_split = 0.2,
    preprocessing_function = None,
    rotation_range = 45,
    zoom_range = 0.2,
    horizontal_flip = True,
    vertical_flip = True,
    fill_mode = 'nearest',
    shear_range = 0.1,
    height_shift_range = 0.1,
    width_shift_range = 0.1
)

train_generator = datagen.flow_from_dataframe(
    dataframe=train,
    directory=data_dir,
    x_col='image_id',
    y_col='Planta',
    target_size=(image_height, image_width),
    batch_size=batch_size,
    class_mode='categorical', # Importante para clasificación multiclase
    subset='training',
    seed=42,
    classes=class_names # Asegurarse del orden de las clases
```

```
)

validation_generator = datagen.flow_from_dataframe(
    dataframe=train,
    directory=data_dir,
    x_col='image_id',
    y_col='Planta',
    target_size=(image_height, image_width),
    batch_size=batch_size,
    class_mode='categorical',
    subset='validation',
    seed=42,
    classes=class_names
)
```

Found 33880 validated image filenames belonging to 9 classes.
 Found 8470 validated image filenames belonging to 9 classes.

5. Cargar el modelo ResNet50 preentrenado (sin la capa superior de clasificación)

```
In [14]: base_model = ResNet50(weights='imagenet', include_top=False, input_shape=(in
```

6. Congelar las capas convolucionales base

```
In [16]: for layer in base_model.layers:
        layer.trainable = False
```

7. Añadir una nueva capa de clasificación personalizada

```
In [18]: x = base_model.output
        x = GlobalAveragePooling2D()(x)
        # Puedes añadir más capas densas aquí si lo deseas
        # x = Dense(1024, activation='relu')(x)
        predictions = Dense(num_classes, activation='softmax')(x)
```

8. Crear el modelo final

```
In [20]: model = Model(inputs=base_model.input, outputs=predictions)
```

9. Compilar el modelo para fine-tuning de la última

capa

```
In [22]: model.compile(optimizer='adam',  
                        loss='categorical_crossentropy',  
                        metrics=['accuracy'])
```

10. Entrenar el modelo con la última capa descongelada

```
In [24]: epochs_fine_tune = 10 # Define el número de épocas para el fine-tuning
```

```
In [25]: print("Comenzando el entrenamiento de la última capa...")  
history_fine_tune = model.fit(  
    train_generator,  
    steps_per_epoch=train_generator.samples // train_generator.batch_size,  
    validation_data=validation_generator,  
    validation_steps=validation_generator.samples // validation_generator.ba  
    epochs=epochs_fine_tune,  
    verbose = 1  
)
```

Comenzando el entrenamiento de la última capa...

Epoch 1/10

2258/2258 ————— **803s** 355ms/step – accuracy: 0.9306 – loss: 0.2436 – val_accuracy: 0.9671 – val_loss: 0.0954

Epoch 2/10

2258/2258 ————— **160s** 71ms/step – accuracy: 1.0000 – loss: 0.0094 – val_accuracy: 0.9637 – val_loss: 0.1062

Epoch 3/10

2258/2258 ————— **798s** 353ms/step – accuracy: 0.9774 – loss: 0.0689 – val_accuracy: 0.9778 – val_loss: 0.0712

Epoch 4/10

2258/2258 ————— **158s** 70ms/step – accuracy: 0.9333 – loss: 0.0986 – val_accuracy: 0.9787 – val_loss: 0.0666

Epoch 5/10

2258/2258 ————— **802s** 355ms/step – accuracy: 0.9823 – loss: 0.0530 – val_accuracy: 0.9856 – val_loss: 0.0491

Epoch 6/10

2258/2258 ————— **160s** 71ms/step – accuracy: 0.9333 – loss: 0.0495 – val_accuracy: 0.9872 – val_loss: 0.0455

Epoch 7/10

2258/2258 ————— **840s** 372ms/step – accuracy: 0.9830 – loss: 0.0544 – val_accuracy: 0.9777 – val_loss: 0.0767

Epoch 8/10

2258/2258 ————— **164s** 73ms/step – accuracy: 1.0000 – loss: 4.3494e-05 – val_accuracy: 0.9800 – val_loss: 0.0726

Epoch 9/10

2258/2258 ————— **823s** 364ms/step – accuracy: 0.9855 – loss: 0.0492 – val_accuracy: 0.9858 – val_loss: 0.0446

Epoch 10/10

2258/2258 ————— **168s** 74ms/step – accuracy: 0.9333 – loss: 0.0542 – val_accuracy: 0.9862 – val_loss: 0.0451

```
In [26]: fin = datetime.datetime.now()
print(f"Fin: {fin.strftime('%Y-%m-%d %H:%M:%S.%f')}")

duracion = fin - inicio
print(f"Duración: {duracion}")
```

Fin: 2025-04-18 20:05:44.010589

Duración: 1:21:17.540637

11. Opcional: Descongelar algunas de las últimas capas del modelo base para un fine-tuning más profundo

```
In [28]: # unfreeze_layers = 5 # Número de capas a descongelar (contando desde la última)
# for layer in base_model.layers[-unfreeze_layers:]
```

```
# layer.trainable = True
```

12. Opcional: Recompilar el modelo después de descongelar las capas

```
In [30]: #model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=1e-5), # Us
#         loss='categorical_crossentropy',
#         metrics=['accuracy'])
```

13. Opcional: Entrenar el modelo con las capas descongeladas

```
In [32]: #epochs_deep_fine_tune = 5 # Define el número de épocas para el fine-tuning

#print(f"\nComenzando el fine-tuning profundo de las últimas {unfreeze_layer}
#history_deep_fine_tune = model.fit(
#    train_generator,
#    steps_per_epoch=train_generator.samples // train_generator.batch_size,
#    validation_data=validation_generator,
#    validation_steps=validation_generator.samples // validation_generator.b
#    epochs=epochs_deep_fine_tune
#)
```

14. Evaluar el modelo final

```
In [34]: print("\nEvaluando el modelo...")
evaluation_result = model.evaluate(validation_generator)
print(f"Pérdida en la validación: {evaluation_result[0]:.4f}")
print(f"Precisión en la validación: {evaluation_result[1]:.4f}")
```

Evaluando el modelo...

565/565  180s 318ms/step - accuracy: 0.9868 - loss: 0.0432

Pérdida en la validación: 0.0386

Precisión en la validación: 0.9880


```
In [50]: # Acceder a los datos de pérdida y precisión
loss = history_fine_tune.history['loss']
val_loss = history_fine_tune.history['val_loss']
accuracy = history_fine_tune.history['accuracy']
val_accuracy = history_fine_tune.history['val_accuracy']

# Graficar la pérdida del modelo
plt.figure(figsize=(12, 4))
```

```
plt.subplot(1, 2, 1)
plt.plot(history_fine_tune.history['loss'], label='Training Loss')
plt.plot(history_fine_tune.history['val_loss'], label='Validation Loss')
plt.title('Loss over epochs')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

# Graficar la precisión del modelo
plt.subplot(1, 2, 2)
plt.plot(history_fine_tune.history['accuracy'], label='Training Accuracy')
plt.plot(history_fine_tune.history['val_accuracy'], label='Validation Accuracy')
plt.title('Accuracy over epochs')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
```

Out[50]: <matplotlib.legend.Legend at 0x28d6dda30>



```
In [52]: epochs = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]
data = {
    'Epoch': epochs,
    'Training Loss': loss,
    'Validation Loss': val_loss,
    'Training Accuracy': accuracy,
    'Validation Accuracy': val_accuracy
}
df = pd.DataFrame(data)

# Imprimir la tabla
print(df)
```

	Epoch	Trainning Loss	Validation Loss	Training Accuracy \
0	1	0.121400	0.095444	0.962203
1	2	0.009383	0.106172	1.000000
2	3	0.066540	0.071185	0.978149
3	4	0.098620	0.066643	0.933333
4	5	0.055798	0.049076	0.981603
5	6	0.049484	0.045450	0.933333
6	7	0.052282	0.076679	0.983730
7	8	0.000043	0.072622	1.000000
8	9	0.049359	0.044633	0.985354
9	10	0.054247	0.045148	0.933333

	Validation Accuracy
0	0.967139
1	0.963712
2	0.977778
3	0.978723
4	0.985579
5	0.987234
6	0.977660
7	0.980024
8	0.985816
9	0.986170

```
In [56]: # Accuracy final en entrenamiento
val_loss = history_fine_tune.history['val_loss'][-1]
```

```
In [58]: print('Perdida final en Train:', history_fine_tune.history['loss'][-1])
print("Pérdida final en Validation:", history_fine_tune.history['val_loss'][-1])

print("Accuracy final en Train:", history_fine_tune.history['accuracy'][-1])
print("Accuracy final en Validation:", history_fine_tune.history['val_accuracy'][-1])
```

```
Perdida final en Train: 0.05424661189317703
Pérdida final en Validation: 0.04514775797724724
Accuracy final en Train: 0.93333333373069763
Accuracy final en Validation: 0.98617023229599
```

```
In [ ]:
```

15. Testeo del Modelo

```
In [66]: test.Planta = test.Planta.astype('str')

test_generator = datagen.flow_from_dataframe(test,
                                             directory = "Datos Finales",
                                             subset = "training",
```


```
x_col = "image_id",  
y_col = "Planta",  
target_size = (image_height, image_width),  
batch_size = batch_size,  
class_mode = "categorical")
```

Found 8472 validated image filenames belonging to 9 classes.

```
In [68]: # -----  
# 1. Obtener predicciones completas para test  
# -----  
y_pred_probs = model.predict(test_generator)  
y_pred = np.argmax(y_pred_probs, axis=1)  
y_true = test_generator.classes
```

565/565 ————— 209s 368ms/step

```
In [69]: # -----  
# 2. Matriz de Confusión  
# -----  
cm = confusion_matrix(y_true, y_pred)  
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=test_gener  
disp.plot(cmap="Blues", xticks_rotation=45)  
plt.title("Matriz de Confusión - Test")  
plt.show()
```



```
In [70]: accuracy = accuracy_score(y_true, y_pred)
print(f"Test Accuracy: {accuracy:.4f}")
```

Test Accuracy: 0.1749

```
In [71]: # Función para normalizar imágenes solo para visualización
def normalize_for_display(img):
    img_min = img.min()
    img_max = img.max()
    if img_max > 1.0 or img_min < 0.0:
        img = (img - img_min) / (img_max - img_min + 1e-7)
    return img

# Obtener nombres de clases
class_indices = test_generator.class_indices
class_names = list(class_indices.keys())

# Reiniciar el generador para empezar desde el principio
test_generator.reset()
```

```
x_batch, y_true_batch = next(test_generator)

# Obtener predicciones del modelo
y_pred_probs_batch = model.predict(x_batch)
y_pred_batch = np.argmax(y_pred_probs_batch, axis=1)
y_true_batch = np.argmax(y_true_batch, axis=1)

# Número de imágenes a mostrar (máximo el tamaño del batch)
num_samples = x_batch.shape[0]

# Mostrar imágenes con etiquetas
plt.figure(figsize=(15, 8))
for i in range(num_samples):
    plt.subplot(3, 4, i + 1)
    img_to_show = normalize_for_display(x_batch[i])
    plt.imshow(img_to_show)
    plt.axis('off')
    plt.title(
        f"Real: {class_names[y_true_batch[i]]}\nPred: {class_names[y_pred_batch[i]]} "
        f"color='green' if y_true_batch[i] == y_pred_batch[i] else 'red' "
    )
plt.tight_layout()
plt.suptitle("Muestreo de Predicciones vs. Etiquetas Reales", fontsize=16, y=1.05)
plt.show()
```

1/1 ————— 1s 791ms/step

```

-----
ValueError                                Traceback (most recent call last)
Cell In[71], line 28
    26 plt.figure(figsize=(15, 8))
    27 for i in range(num_samples):
--> 28     plt.subplot(3, 4, i + 1)
    29     img_to_show = normalize_for_display(x_batch[i])
    30     plt.imshow(img_to_show)

File /opt/anaconda3/lib/python3.12/site-packages/matplotlib/pyplot.py:1534,
in subplot(*args, **kwargs)
    1531 fig = gcf()
    1533 # First, search for an existing subplot with a matching spec.
-> 1534 key = SubplotSpec._from_subplot_args(fig, args)
    1536 for ax in fig.axes:
    1537     # If we found an Axes at the position, we can reuse it if the us
er passed no
    1538     # kwargs or if the Axes class and kwargs are identical.
    1539     if (ax.get_subplotspec() == key
    1540         and (kwargs == {}
    1541             or (ax._projection_init
    1542                 == fig._process_projection_requirements(**kwargs)))
    ):

File /opt/anaconda3/lib/python3.12/site-packages/matplotlib/gridspec.py:589,
in SubplotSpec._from_subplot_args(figure, args)
    587 else:
    588     if not isinstance(num, Integral) or num < 1 or num > rows*cols:
--> 589         raise ValueError(
    590             f"num must be an integer with 1 <= num <= {rows*cols}, "
    591             f"not {num!r}"
    592         )
    593     i = j = num
    594 return gs[i-1:j]

ValueError: num must be an integer with 1 <= num <= 12, not 13

```



```
In [76]: # 3. Convertir etiquetas reales a one-hot
n_classes = y_pred_probs.shape[1]
y_true_onehot = to_categorical(y_true, num_classes=n_classes)

# 4. Calcular ROC y AUC por clase
fpr = dict()
tpr = dict()
roc_auc = dict()


for i in range(n_classes):
    fpr[i], tpr[i], _ = roc_curve(y_true_onehot[:, i], y_pred_probs[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])

# 5. Obtener nombres de clase
class_names = list(test_generator.class_indices.keys())

# 6. Graficar curvas ROC
plt.figure(figsize=(12, 8))
for i in range(n_classes):
    plt.plot(fpr[i], tpr[i], label=f'{class_names[i]} (AUC = {roc_auc[i]:.2f})')

plt.plot([0, 1], [0, 1], 'k--', label='Aleatorio')
plt.xlabel('Tasa de Falsos Positivos')
plt.ylabel('Tasa de Verdaderos Positivos')
plt.title('Curvas ROC por clase')
plt.legend(loc='lower right')
plt.grid(True)
```

```
plt.show()
```



```
In [78]: from sklearn.metrics import roc_auc_score

# AUC macro (promedio de clases, todas igual de importantes)
auc_macro = roc_auc_score(y_true_onehot, y_pred_probs, average='macro')

# AUC micro (promedia sobre todas las instancias)
auc_micro = roc_auc_score(y_true_onehot, y_pred_probs, average='micro')

print(f"AUC Macro promedio: {auc_macro:.4f}")
print(f"AUC Micro promedio: {auc_micro:.4f}")
```

AUC Macro promedio: 0.5053
AUC Micro promedio: 0.5511

```
In [80]: # Asegúrate de tener esto:
# y_pred = np.argmax(y_pred_probs, axis=1)
# y_true = test_generator.classes

# Accuracy
accuracy = accuracy_score(y_true, y_pred)

# F1 Scores
f1_macro = f1_score(y_true, y_pred, average='macro')
```

```

f1_micro = f1_score(y_true, y_pred, average='micro')
f1_weighted = f1_score(y_true, y_pred, average='weighted')

print(f"Accuracy: {accuracy:.4f}")
print(f"F1 Score (macro): {f1_macro:.4f}")
print(f"F1 Score (micro): {f1_micro:.4f}")
print(f"F1 Score (weighted): {f1_weighted:.4f}")

# Detalle por clase
class_names = list(test_generator.class_indices.keys())
print("\nReporte de clasificación por clase:")
print(classification_report(y_true, y_pred, target_names=class_names))

```

Accuracy: 0.1749
 F1 Score (macro): 0.0796
 F1 Score (micro): 0.1749
 F1 Score (weighted): 0.1742

Reporte de clasificación por clase:

	precision	recall	f1-score	support
Apple	0.06	0.40	0.11	523
Corn_(maize)	0.07	0.03	0.04	623
Potato	0.03	0.03	0.03	337
Strawberry	0.02	0.06	0.03	245
Tomato	0.34	0.35	0.35	2914
Yuca	0.42	0.06	0.11	3418
bean	0.01	0.00	0.01	208
cauliflower	0.03	0.03	0.03	95
cucumber	0.01	0.03	0.02	109
accuracy			0.17	8472
macro avg	0.11	0.11	0.08	8472
weighted avg	0.30	0.17	0.17	8472

In []: